

# Training a Neural Network

Michael Shao

# Machine Learning

- Looking for the function

$$f(x_0, x_1, \dots) = \hat{y}$$

# We've defined the structure

- but what about the parameters?

$f_{\theta}$

$w_0, w_1, w_2, \dots$

$b_0, b_1, b_2, \dots$

$\theta$

# What's the best $\theta$ ?

- We can evaluate the performance of the model based on how closely the output follows the distribution in the **labeled** training set.



$x$



$y$

# Comparing distributions?

- Defining a loss function

$L(\theta) = \text{Difference}$

be

$f_{\theta}($



$x)$  and



$y$

# Optimization

Gradient Descent

Minimizing the loss function

$$\theta^{\star} = \arg \min_{\theta} L(\theta)$$

How do we do this?


# Gradient Descent

- Descending a slope:


$$\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta)$$

<https://uclaacm.github.io/gradient-descent-visualiser/#playground>







# Local Minima



# Convergence



# Divergence



# Step, Batch, Epoch

- Split training data randomly into batches based on batch size.
- Calculate loss on batch, perform gradient descent = 1 step
- Iterate over all batches = 1 epoch

# Loss Functions

The Training Criteria

# Mean Squared Error Loss

`criterion = torch.nn.MSELoss()`

$$\frac{1}{N} \sum_{i=1}^N (f_{\theta}(x_i) - y_i)^2$$



# Loss for Classification Tasks?

- Number of misclassifications?
- Inaccuracy?



# What does a classification model output?

0	0.2
1	1.3
2	(25.7)
3	−11
4	102

But each output has no “meaning”. Their values are only relative to the oth



# Introducing... Softmax Activation

- Map values to (0, 1) with a sum of 1.

$$\sigma \begin{pmatrix} 0.2 \\ 1.3 \\ 25.7 \\ -11 \\ 102 \end{pmatrix} = \begin{pmatrix} 6.15 \times 10^{-45} \\ 1.85 \times 10^{-44} \\ 7.3 \times 10^{-34} \\ 8.41 \times 10^{-50} \\ 1 \end{pmatrix}$$

Now, these values can be interpreted as probabilities:  $\begin{pmatrix} p_0 \\ p_1 \\ \dots \\ p_C \end{pmatrix}$ .

# Cross Entropy Loss

- We want the corresponding probability to be close to 1, so...
- We want the log probability to be close to 0. In other words,
- We want to minimize  $-\log(p_y)$
- $L(\theta) = -\frac{1}{N} \sum_{i=1}^N \log(p_{i,y_i})$

# Cross Entropy Loss

- `criterion = torch.nn.CrossEntropyLoss()`
- `criterion = torch.nn.BCELoss()` # For Binary CE

# Calculating The Gradient

Back Propagation... In very very very simple words

# Back Propagation

- The gradient is propagated from the last layer all the way back to the first layer.
- Keywords: partial derivatives, chain rule
- <https://www.youtube.com/watch?v=llg3gGewQ5U>